



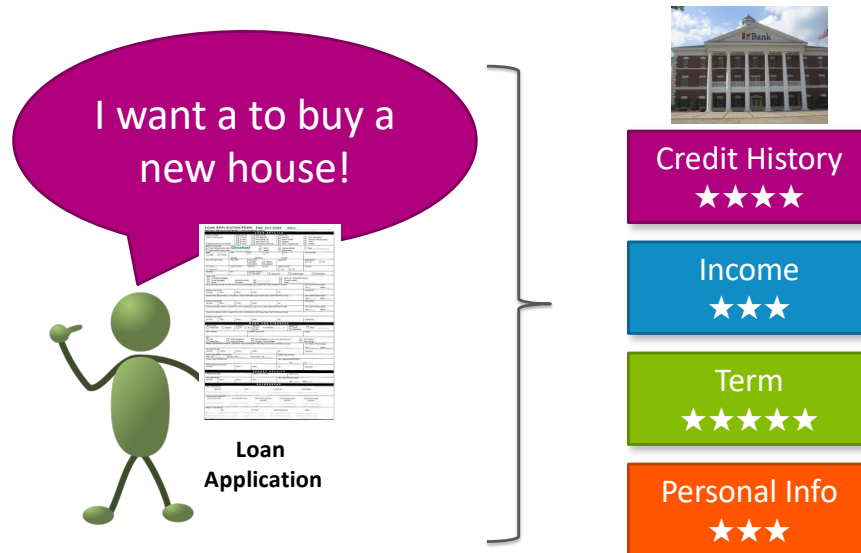
Decision Trees

CS229: Machine Learning
Carlos Guestrin
Stanford University

Slides include content developed by and co-developed with Emily Fox

Predicting potential loan defaults

What makes a loan risky?




Credit history explained

Did I pay previous
loans on time?




Example: excellent,
good, or fair



Income

What's my income?


Example:
\$80K per year



Loan terms

How soon do I need to
pay the loan?


Example: 3 years,
5 years,...




Personal information

Age, reason for the loan,
marital status,...

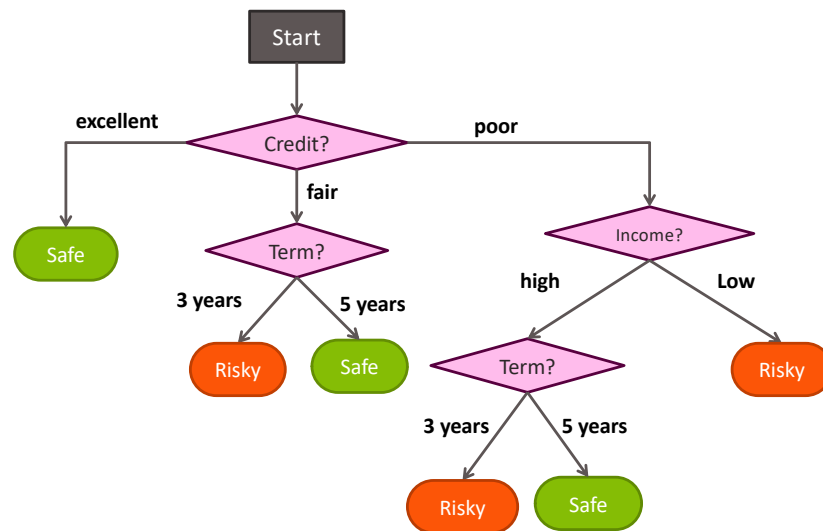
Example: Home loan for a
married couple



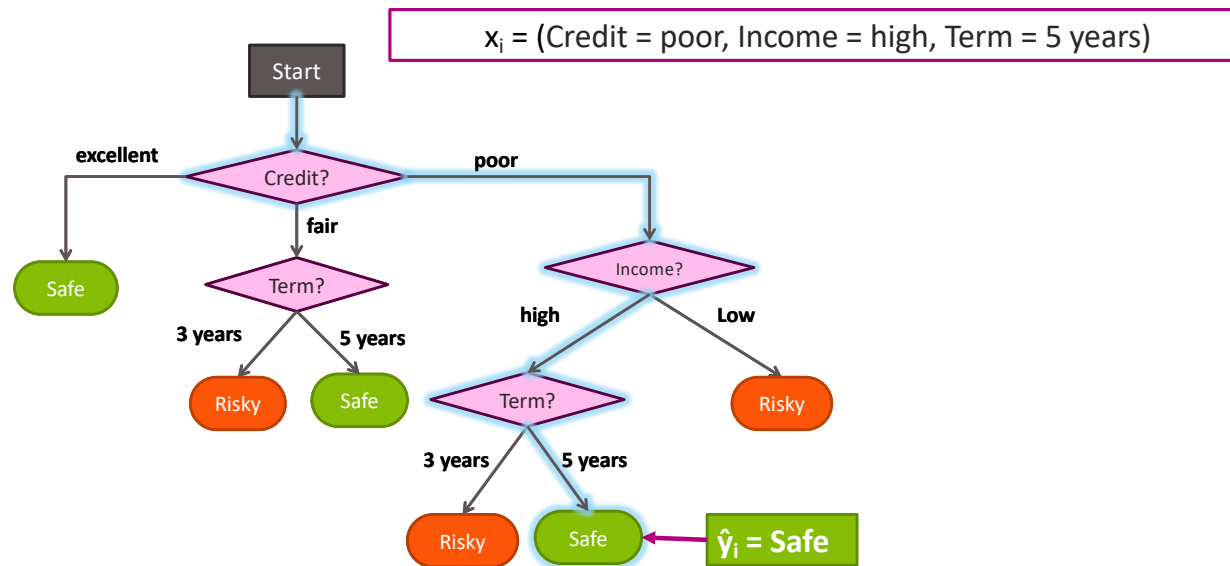
Classifier review



This module ... decision trees



Scoring a loan application




Decision tree learning task

Decision tree learning problem

Training data: N observations (x_i, y_i)

| Credit | Term | Income | y |
|-----------|-------|--------|-------|
| excellent | 3 yrs | high | safe |
| fair | 5 yrs | low | risky |
| fair | 3 yrs | high | safe |
| poor | 5 yrs | high | risky |
| excellent | 3 yrs | low | risky |
| fair | 5 yrs | low | safe |
| poor | 3 yrs | high | risky |
| poor | 5 yrs | low | safe |
| fair | 3 yrs | high | safe |



Quality metric: Classification error

- Error measures fraction of mistakes


$$\text{Error} = \frac{\text{\# incorrect predictions}}{\text{\# examples}}$$

- Best possible value : 0.0
- Worst possible value: 1.0

How do we find the best tree?

Exponentially large number of possible trees
makes decision tree learning **hard**!

Learning the smallest
decision tree is an
NP-hard problem
[Hyafil & Rivest '76]



Greedy decision tree learning


Our training data table

Assume $N = 40$, 3 features


| Credit | Term | Income | y |
|-----------|-------|--------|-------|
| excellent | 3 yrs | high | safe |
| fair | 5 yrs | low | risky |
| fair | 3 yrs | high | safe |
| poor | 5 yrs | high | risky |
| excellent | 3 yrs | low | risky |
| fair | 5 yrs | low | safe |
| poor | 3 yrs | high | risky |
| poor | 5 yrs | low | safe |
| fair | 3 yrs | high | safe |

Start with all the data

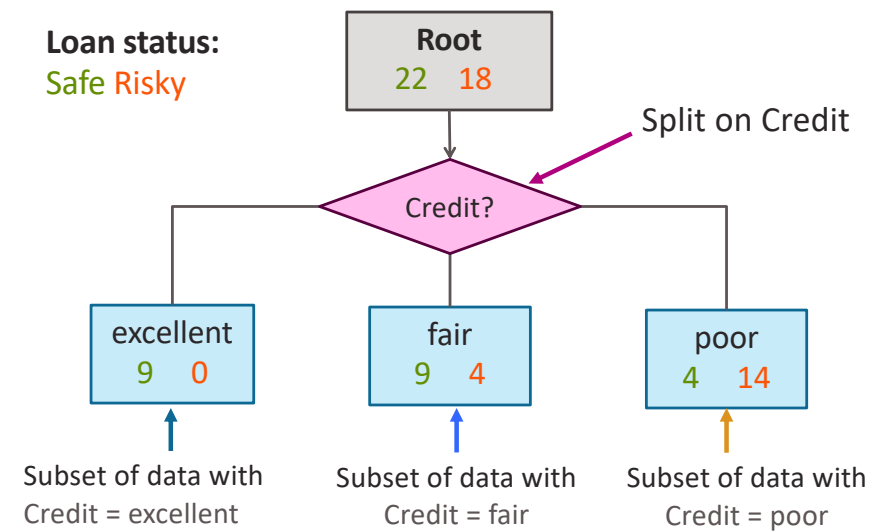
Loan status: **Safe** **Risky**



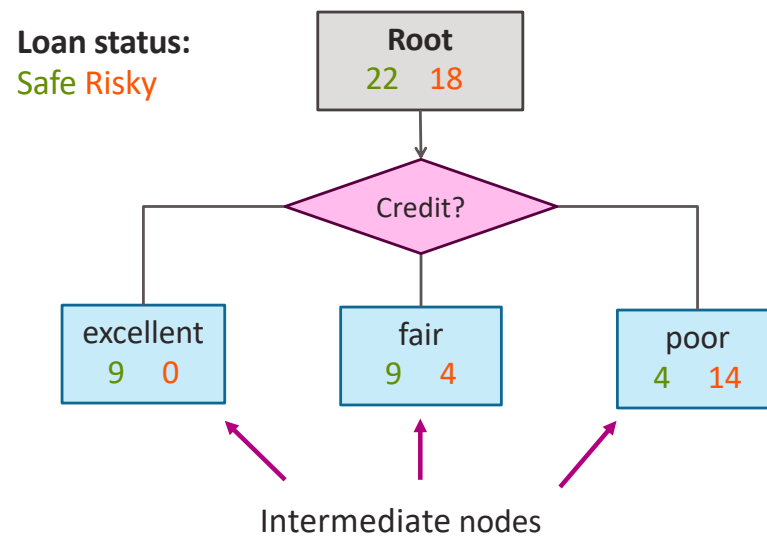
Compact visual notation: Root node



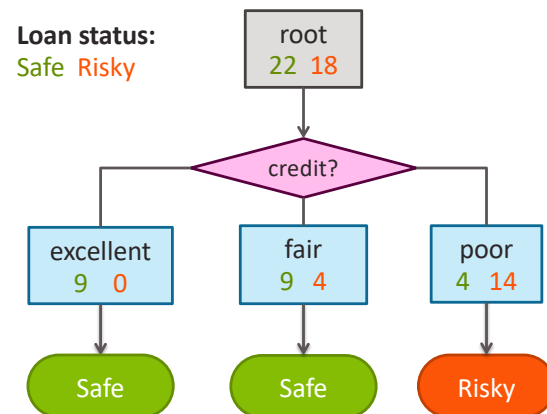
Decision stump: Single level tree



Visual notation: Intermediate nodes



Making predictions with a decision stump



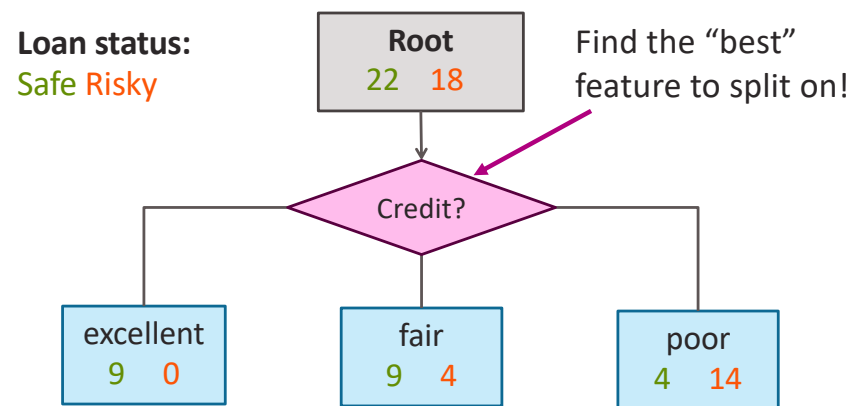
For each intermediate node,
set \hat{y} = majority value



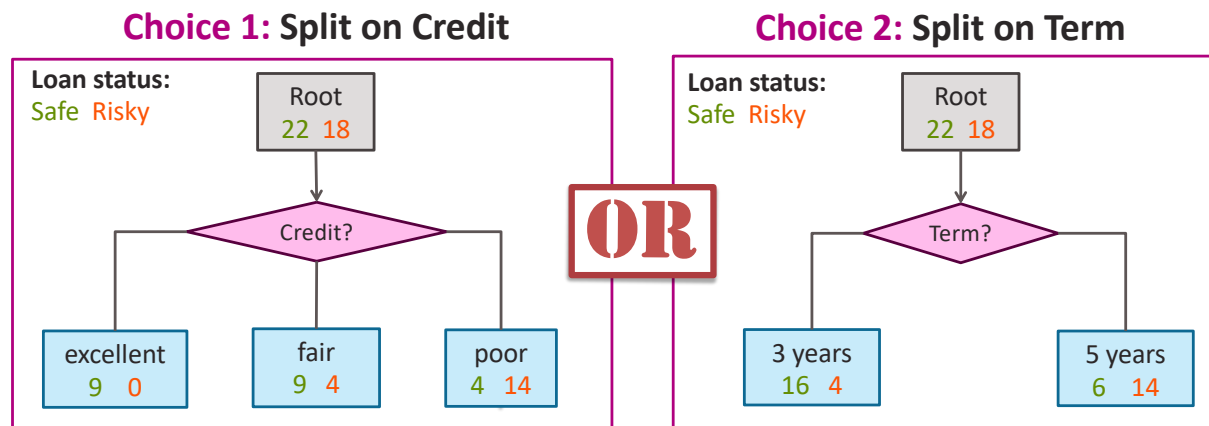
Selecting best feature to split on



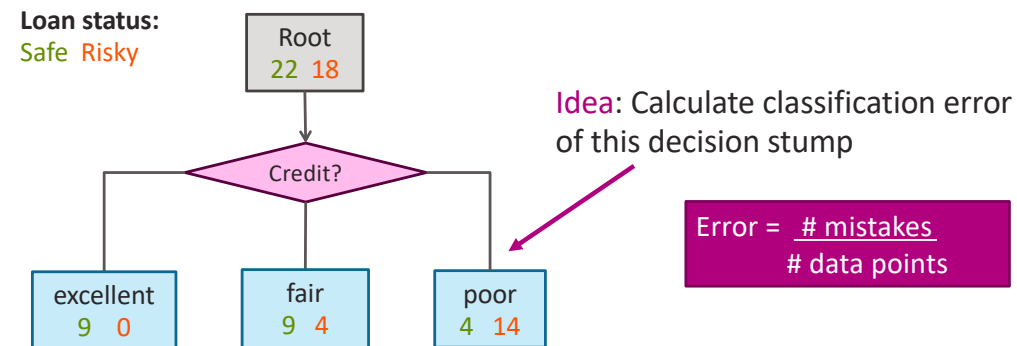
How do we learn a decision stump?



How do we select the best feature?




How do we measure effectiveness of a split?



Calculating classification error

- **Step 1:** \hat{y} = class of majority of data in node
- **Step 2:** Calculate classification error of predicting \hat{y} for this data

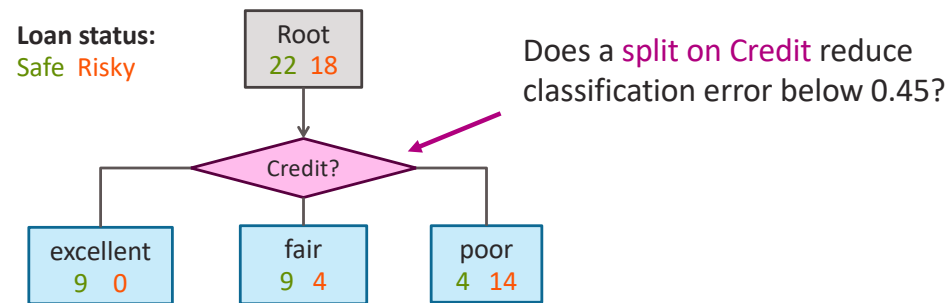


Error = _____
=

| Tree | Classification error |
|--------|----------------------|
| (root) | 0.45 |

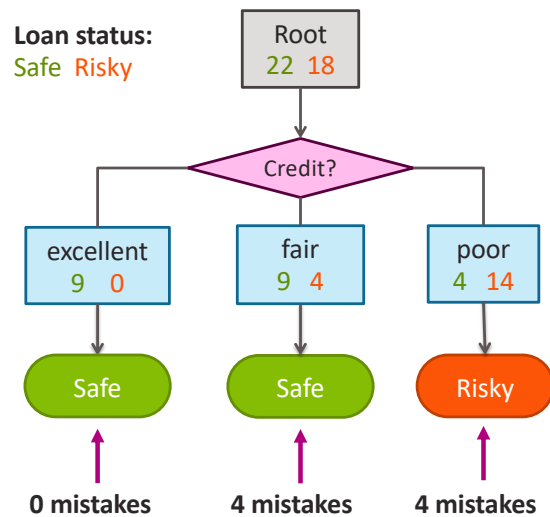
Choice 1: Split on **Credit** history?

Choice 1: Split on Credit



Split on Credit: Classification error

Choice 1: Split on Credit

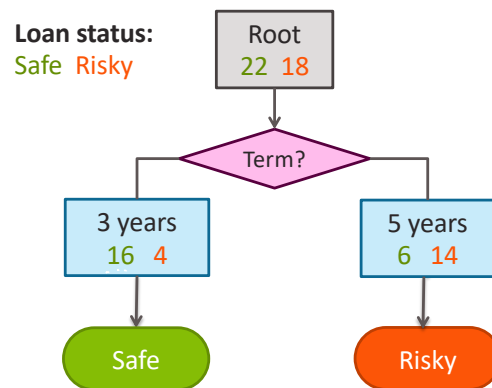


Error = _____
=

| Tree | Classification error |
|-----------------|----------------------|
| (root) | 0.45 |
| Split on credit | 0.2 |


Choice 2: Split on Term?

Choice 2: Split on Term



Evaluating the split on Term

Choice 2: Split on Term



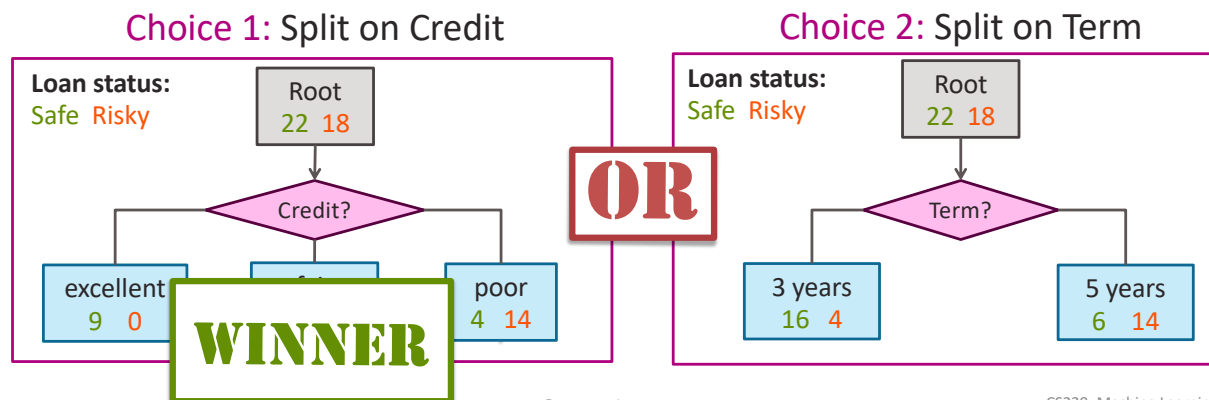
Error = _____

=

| Tree | Classification error |
|-----------------|----------------------|
| (root) | 0.45 |
| Split on credit | 0.2 |
| Split on term | 0.25 |

Choice 1 vs Choice 2: Comparing split on Credit vs Term

| Tree | Classification error |
|--------------------|----------------------|
| (root) | 0.45 |
| split on credit | 0.2 |
| split on loan term | 0.25 |



Feature split selection algorithm

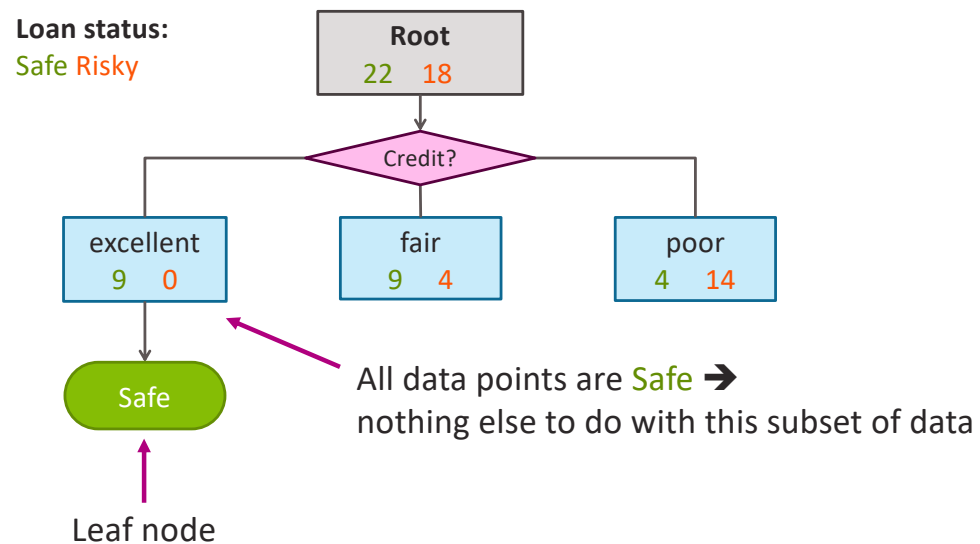
- Given a subset of data M (a node in a tree)
- For each feature $h_i(x)$:
 1. Split data of M according to feature $h_i(x)$
 2. Compute classification error of split
- Chose feature $h^*(x)$ with lowest classification error



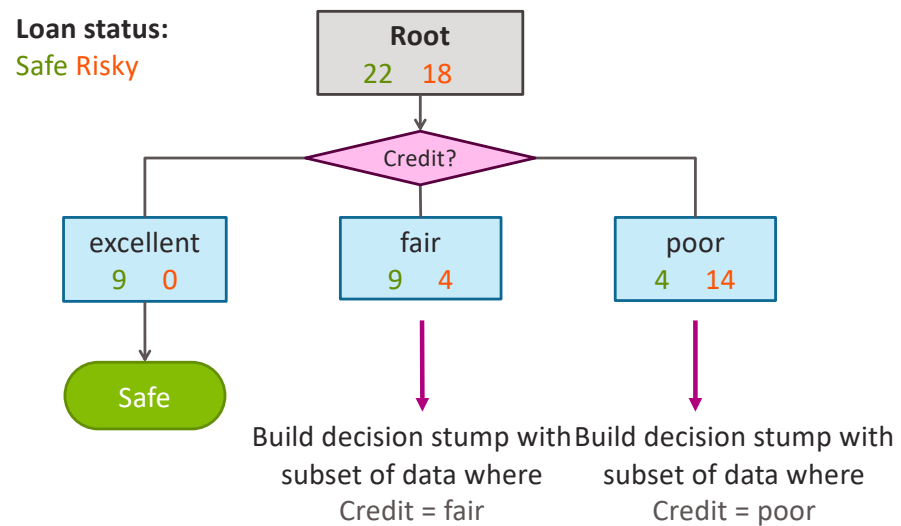
Recursion & Stopping conditions



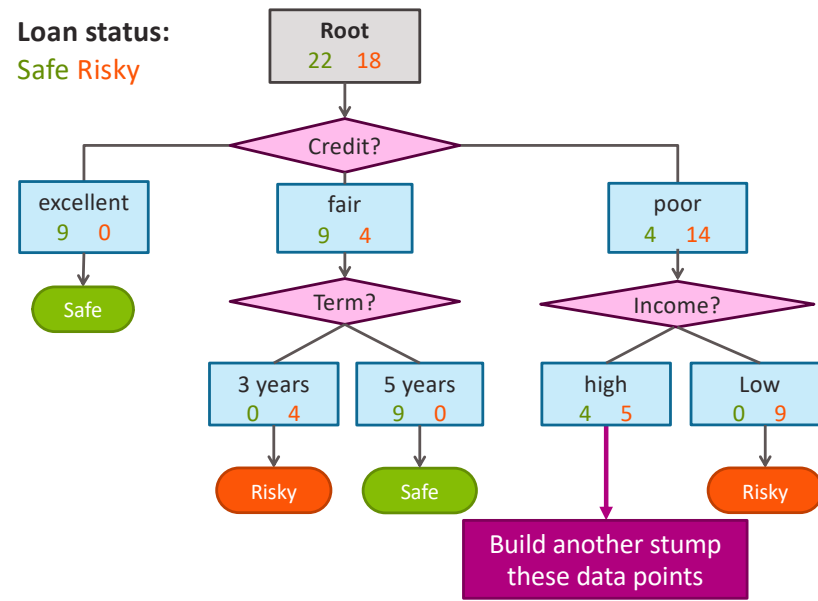
We've learned a decision stump, what next?



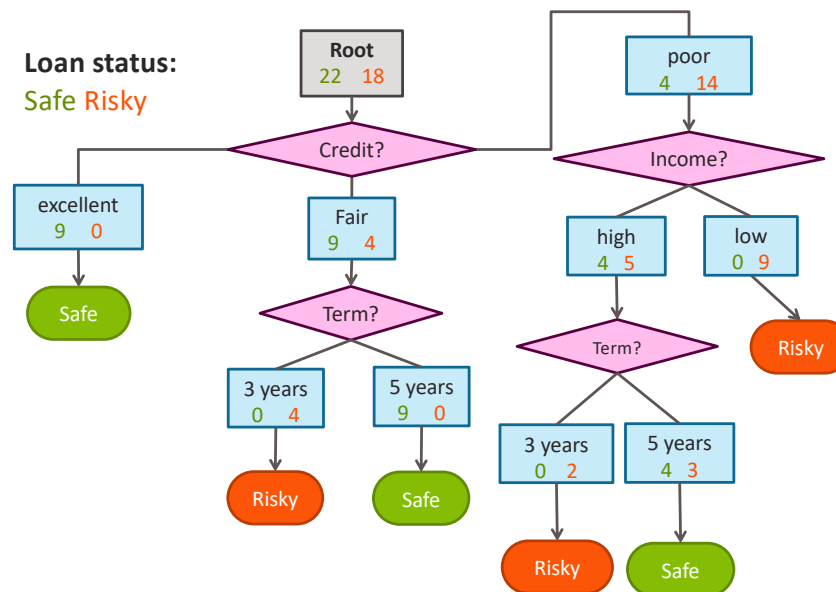
Tree learning = Recursive stump learning



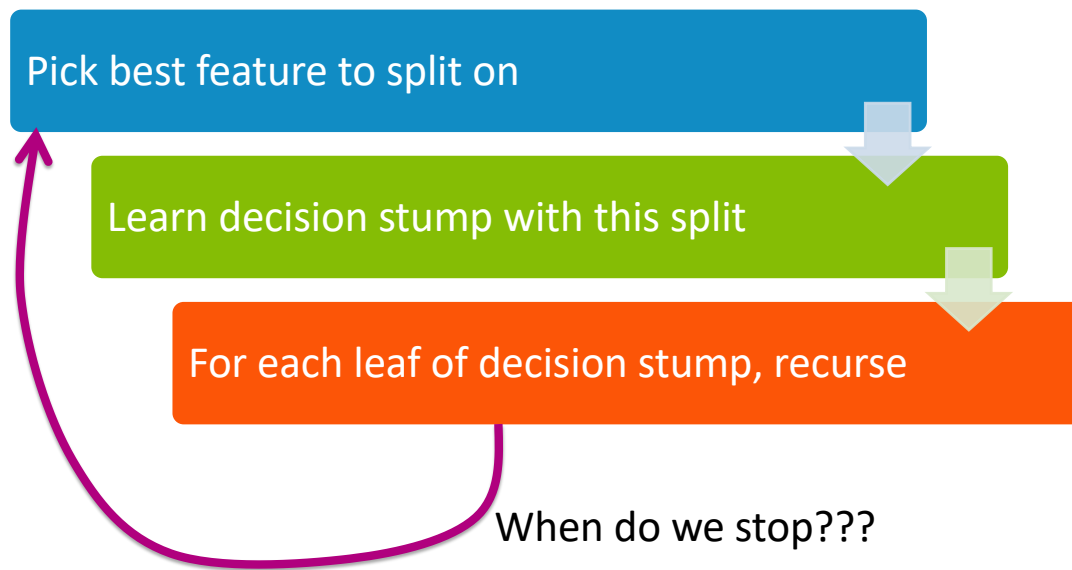
Second level



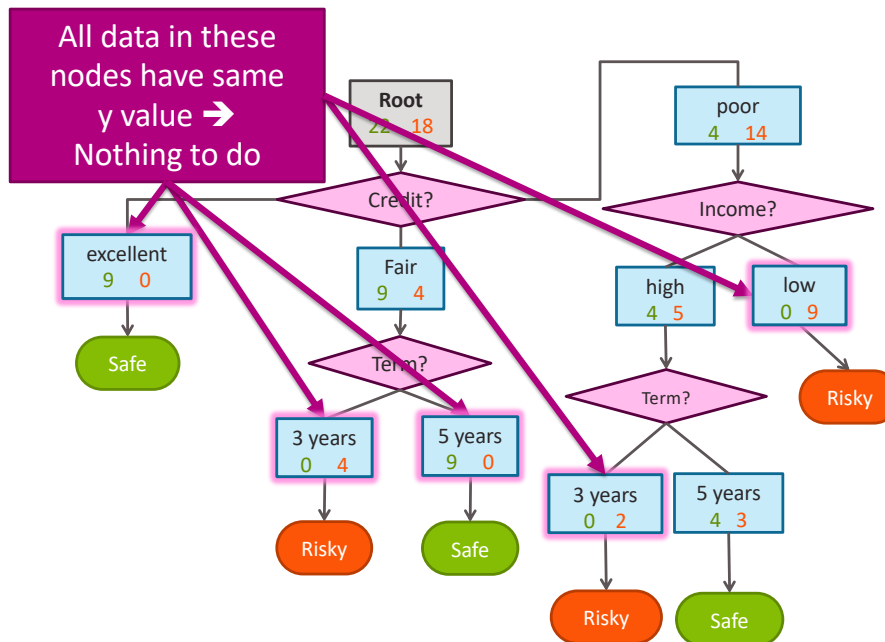
Final decision tree



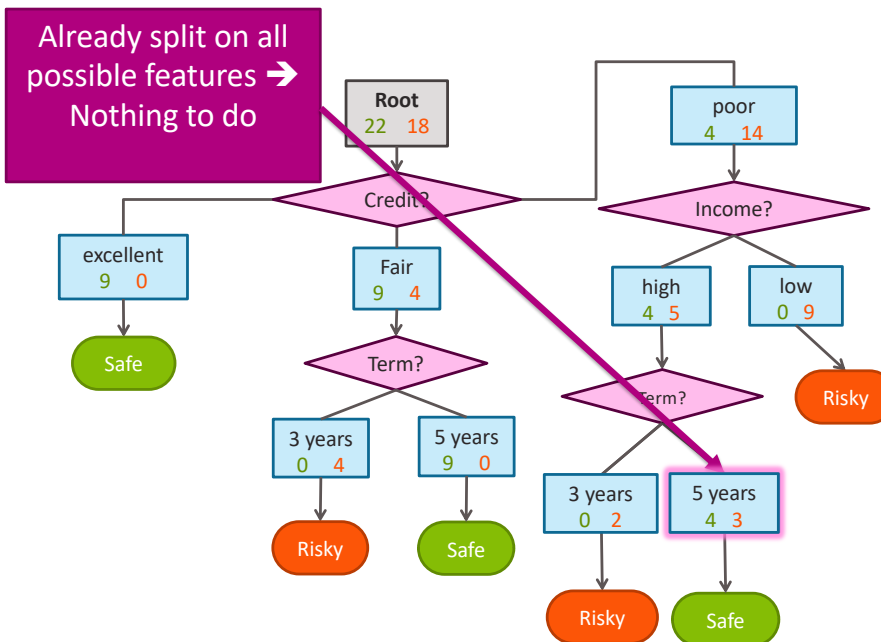
Simple greedy decision tree learning



Stopping condition 1: All data agrees on y



Stopping condition 2: Already split on all features



Greedy decision tree learning

- **Step 1:** Start with an empty tree

- **Step 2:** Select a feature to split data

- For each split of the tree:

- **Step 3:** If nothing more to do,
make predictions

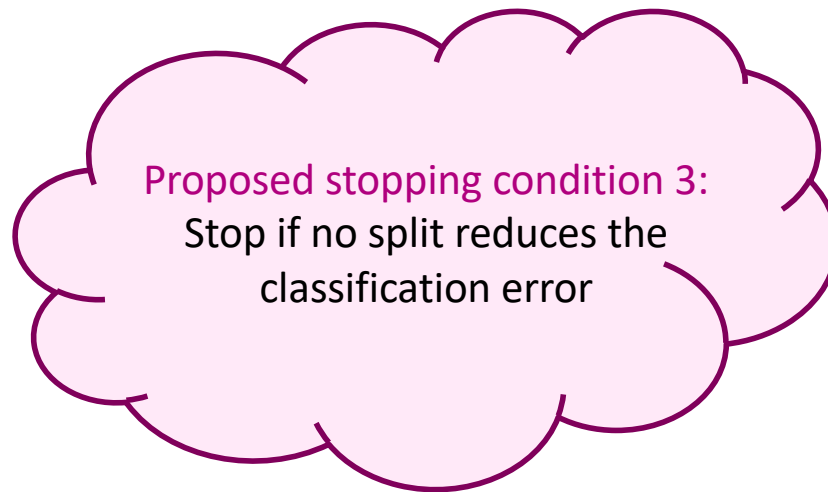
- **Step 4:** Otherwise, go to **Step 2** &
continue (recurse) on this split

Pick feature split
leading to lowest
classification error

Stopping conditions
1 & 2

Recursion

Is this a good idea?



Stopping condition 3:
Don't stop if error doesn't decrease???

$$y = x[1] \text{ xor } x[2]$$

| x[1] | x[2] | y |
|-------|-------|-------|
| False | False | False |
| False | True | True |
| True | False | True |
| True | True | False |

y values
True False

| Root |
|------|
| 2 2 |

Error = _____
=


| Tree | Classification error |
|--------|----------------------|
| (root) | 0.5 |

Consider split on $x[1]$

$$y = x[1] \text{ xor } x[2]$$

| $x[1]$ | $x[2]$ | y |
|--------|--------|-------|
| False | False | False |
| False | True | True |
| True | False | True |
| True | True | False |

y values
 True False



Error = _____
 =


| Tree | Classification error |
|-----------------|----------------------|
| (root) | 0.5 |
| Split on $x[1]$ | 0.5 |

Consider split on $x[2]$

$$y = x[1] \text{ xor } x[2]$$

| $x[1]$ | $x[2]$ | y |
|--------|--------|-------|
| False | False | False |
| False | True | True |
| True | False | True |
| True | True | False |

y values
True False



$$\text{Error} = \frac{1+1}{2+2} = 0.5$$

Neither features
improve training error...
Stop now???


| Tree | Classification error |
|-----------------|----------------------|
| (root) | 0.5 |
| Split on $x[1]$ | 0.5 |
| Split on $x[2]$ | 0.5 |

Final tree with stopping condition 3

$$y = x[1] \text{ xor } x[2]$$

| x[1] | x[2] | y |
|-------|-------|-------|
| False | False | False |
| False | True | True |
| True | False | True |
| True | True | False |

y values
True False



| Tree | Classification error |
|---------------------------|----------------------|
| with stopping condition 3 | 0.5 |

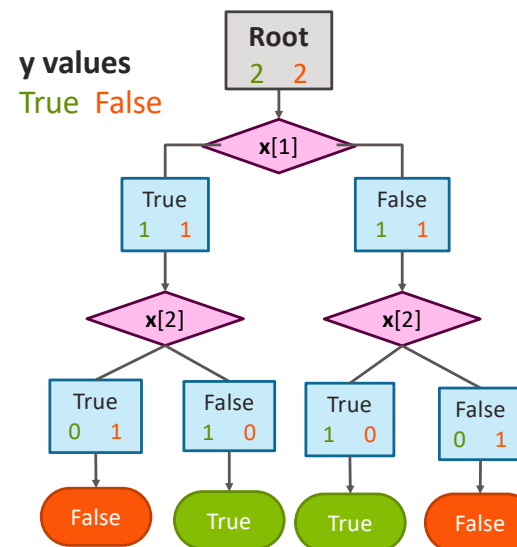
Without stopping condition 3

Condition 3 (stopping when training error doesn't improve) is not recommended!

$$y = x[1] \text{ xor } x[2]$$

| x[1] | x[2] | y |
|-------|-------|-------|
| False | False | False |
| False | True | True |
| True | False | True |
| True | True | False |

| Tree | Classification error |
|------------------------------|----------------------|
| with stopping condition 3 | 0.5 |
| without stopping condition 3 | |

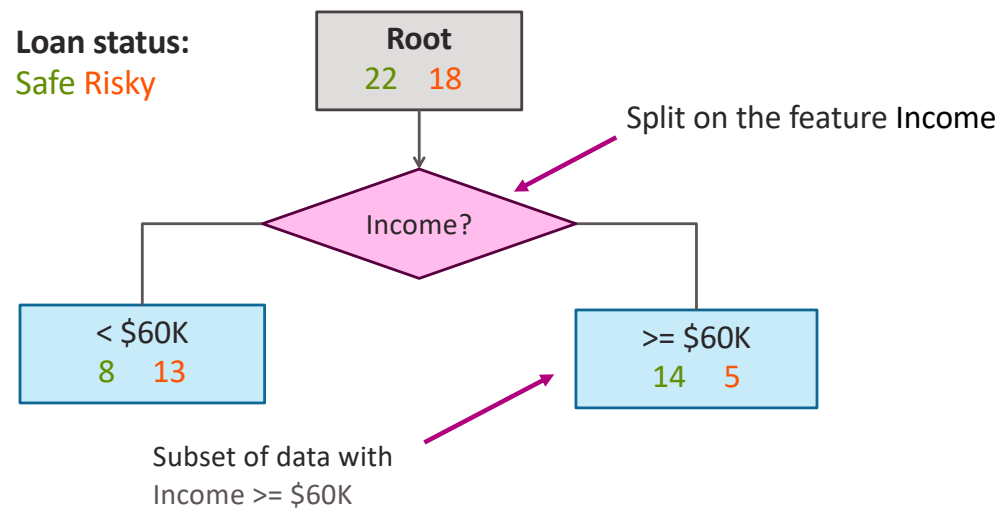


Decision tree learning:
Real valued features

How do we use real values inputs?

| Income | Credit | Term | y |
|---------|-----------|-------|-------|
| \$105 K | excellent | 3 yrs | Safe |
| \$112 K | good | 5 yrs | Risky |
| \$73 K | fair | 3 yrs | Safe |
| \$69 K | excellent | 5 yrs | Safe |
| \$217 K | excellent | 3 yrs | Risky |
| \$120 K | good | 5 yrs | Safe |
| \$64 K | fair | 3 yrs | Risky |
| \$340 K | excellent | 5 yrs | Safe |
| \$60 K | good | 3 yrs | Risky |


Threshold split



Finding the best threshold split


Infinite possible
values of t

Income = t^*




Consider a threshold between points

Same **classification error** for any threshold split between v_A and v_B




Only need to consider mid-points




Threshold split selection algorithm

- **Step 1:** Sort the values of a feature $h_j(x)$:
Let $\{v_1, v_2, v_3, \dots, v_N\}$ denote sorted values
- **Step 2:**
 - For $i = 1 \dots N-1$
 - Consider split $t_i = (v_i + v_{i+1}) / 2$
 - Compute classification error for threshold split $h_j(x) \geq t_i$
 - Chose the t^* with the lowest classification error


Visualizing the threshold split




Split on Age ≥ 38



Depth 2: Split on Income \geq \$60K




Each split partitions the 2-D space





Decision trees vs logistic regression: *Example*

Logistic regression

| Feature | Value | Weight Learned |
|----------|--------|----------------|
| $h_0(x)$ | 1 | 0.22 |
| $h_1(x)$ | $x[1]$ | 1.12 |
| $h_2(x)$ | $x[2]$ | -1.07 |




Depth 1: Split on $x[1]$




62

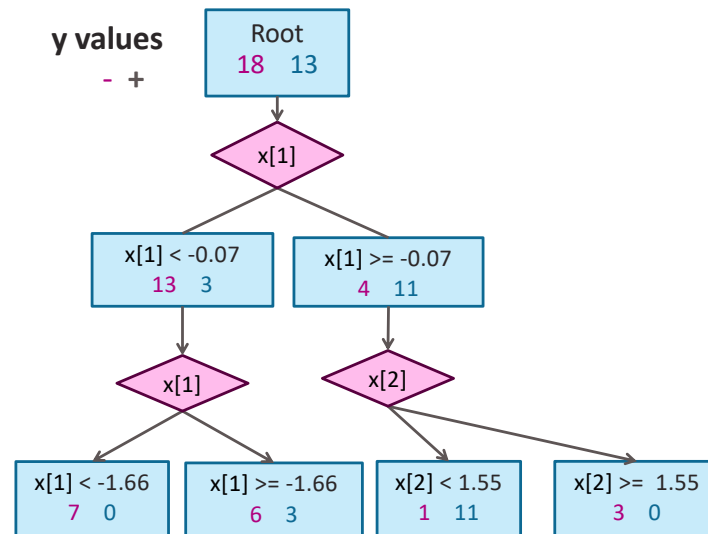
y values
- +



Depth 2



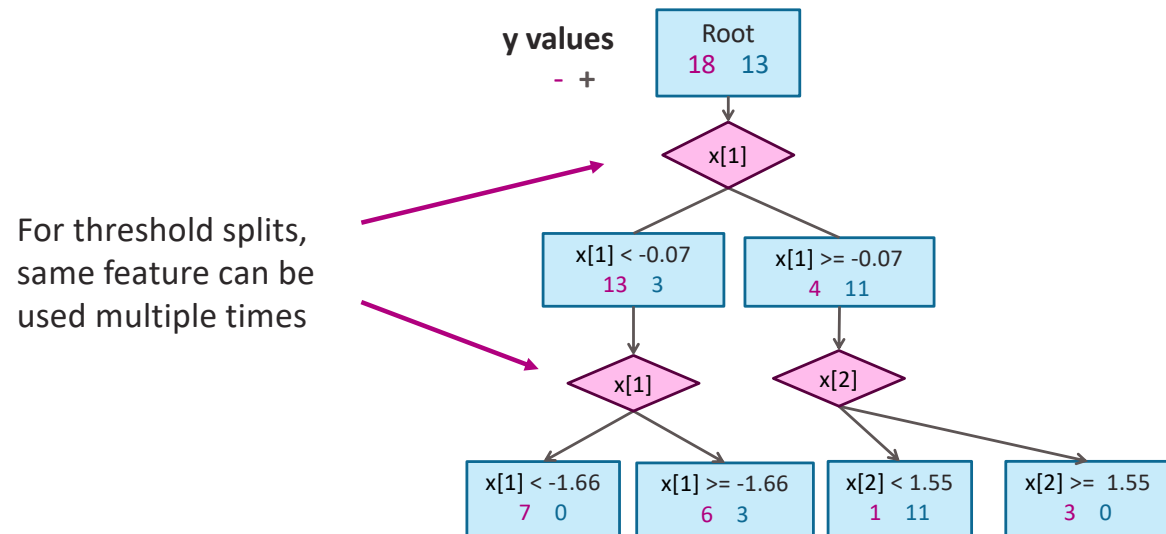
63







©2021 Carlos Guestrin

CS229: Machine Learning

Threshold split caveat




Decision boundaries




Comparing decision boundaries

Decision Tree



Logistic Regression



Summary of decision trees

What you can do now

- Define a decision tree classifier
- Interpret the output of a decision trees
- Learn a decision tree classifier using greedy algorithm
- Traverse a decision tree to make predictions
 - Majority class predictions
- Tackle continuous and discrete features